

einsum_decomposition

March 10, 2022

1 Einsum decomposition

This notebook shows a way to decompose `einsum` into a subset of operations (`expand_dims`, `squeeze`, `transpose`, extended matrix multiplication).

```
[1]: from jyquickhelper import add_notebook_menu
      add_notebook_menu()
```

```
[1]: <IPython.core.display.HTML object>
```

```
[2]: %load_ext mlproduct
```

1.1 Operator explanation with equation `bac,cd,def=ebc`

The operator `einsum` takes an equation and some inputs. Every letter involved in the equation is a loop. Let's see on one example.

```
[3]: import numpy

m1 = numpy.arange(0, 8).astype(numpy.float32).reshape((2, 2, 2)) + 10
m2 = numpy.arange(0, 4).astype(numpy.float32).reshape((2, 2)) + 100
m3 = numpy.arange(0, 8).astype(numpy.float32).reshape((2, 2, 2)) + 1000

equation = "bac,cd,def->ebc"
truth = numpy.einsum(equation, m1, m2, m3)
truth
```

```
[3]: array([[[ 8866198.,  9864696.],
             [12090270., 13152928.]],

           [[ 8883886.,  9884376.],
             [12114390., 13179168.]]], dtype=float32)
```

This summation is equivalent to:

```
[4]: res = numpy.zeros((2, 2, 2))
      for a in range(0, 2):
          for b in range(0, 2):
              for c in range(0, 2):
                  for d in range(0, 2):
                      for e in range(0, 2):
                          for f in range(0, 2):
                              res[e, b, c] += m1[b, a, c] * m2[c, d] * m3[d, e, f]
```

```
res
```

```
[4]: array([[[ 8866198.,  9864696.],
             [12090270., 13152928.]],

           [[ 8883886.,  9884376.],
             [12114390., 13179168.]])
```

Theoretically, this summation in this case has a cost of $O(N^6)$. However this simple computation is usually much longer than using matrix multiplications along the path. $O(N^4)$ is the cost of the heaviest matrix multiplication in this case). But to do that, the equation needs to be decomposed into a sequence of matrix multiplications.

1.1.1 Decomposition of `bac,cd,def=ebc`

```
[5]: import numpy
      from mlproduct.testing.einsum import (
          decompose_einsum_equation, apply_einsum_sequence)
```

```
[6]: m1 = numpy.arange(0, 8).astype(numpy.float32).reshape((2, 2, 2)) + 10
      m2 = numpy.arange(0, 4).astype(numpy.float32).reshape((2, 2)) + 100
      m3 = numpy.arange(0, 8).astype(numpy.float32).reshape((2, 2, 2)) + 1000
```

```
[7]: seq = decompose_einsum_equation("bac,cd,def->ebc")
```

```
[8]: from jyquickhelper import RenderJsDot
      RenderJsDot(seq.to_dot(size=7))
```

```
[8]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x1c6d45f6250>
```

Then the result can be obtained as follows:

```
[9]: apply_einsum_sequence(seq, m1, m2, m3)
```

```
[9]: array([[[ 8866198.,  9864696.],
             [12090270., 13152928.]],

           [[ 8883886.,  9884376.],
             [12114390., 13179168.]]) dtype=float32)
```

1.1.2 operator `matmul`

This operator can be used to represent either a multiplication, either a matrix multiplication but it applies only on arrays with the same number of dimensions. It can be broken into multiplication of matrix multiplication.

```
[10]: seq_clean = decompose_einsum_equation("bac,cd,def->ebc", strategy='numpy', clean=True)
      RenderJsDot(seq_clean.to_dot(size=7))
```

```
[10]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x1c6d45f6910>
```

Operator `transpose_mm` is a regular transposition, it takes two inputs but only transposes the first input before returning it. Operator `batch_dot` is a matrix multiplication. It is left that way on purpose as it may be implemented with function `dot` or `gemm`. The operator distinguishes between 3 kind of axes: batch axes, kept axes, sum(mation) axes. It then reshapes both input matrices with 3D tensors, batch axis, row axis, column axis to use function `numpy.dot`.

1.1.3 ONNX

The previous graph can be converted into ONNX.

```
[11]: onx = seq_clean.to_onnx("Y", "X1", "X2", "X3", dtype=numpy.float32)
# with open("einsum.onnx", "wb") as f:
#     f.write(onx.SerializeToString())
%onnxview onx
```

```
[11]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x1c6d4631640>
```

```
[12]: from onnxruntime import InferenceSession
sess = InferenceSession(onx.SerializeToString())
sess.run(None, {'X1': m1.astype(numpy.float32),
                'X2': m2.astype(numpy.float32),
                'X3': m3.astype(numpy.float32)})[0]
```

```
[12]: array([[ 8866198.,  9864696.],
            [12090270., 13152928.]],

           [[ 8883886.,  9884376.],
            [12114390., 13179168.]]], dtype=float32)
```

1.1.4 onnxruntime

```
[13]: import onnx
from onnx import helper, numpy_helper
from onnxruntime import InferenceSession

def make_model1(equation):
    model = helper.make_model(
        opset_imports=[helper.make_operatorsetid('', 13)],
        graph=helper.make_graph(
            name='einsum_test',
            inputs=[helper.make_tensor_value_info("X", onnx.TensorProto.FLOAT, None),
                    helper.make_tensor_value_info("Y", onnx.TensorProto.FLOAT, None),
                    helper.make_tensor_value_info("Z", onnx.TensorProto.FLOAT, None)],
            outputs=[helper.make_tensor_value_info("A", onnx.TensorProto.FLOAT, None)],
            nodes=[
                helper.make_node("Einsum", ["X", "Y", "Z"], ["A"], equation=equation)
            ]
        )
    )
    return model

model = make_model1("bac,cd,def->ebc")
sess = InferenceSession(model.SerializeToString())
```

```
[14]: sess.run(None, {'X': m1.astype(numpy.float32),
                    'Y': m2.astype(numpy.float32),
                    'Z': m3.astype(numpy.float32)})[0]
```

```
[14]: array([[[ 8866198.,  9864696.],
              [12090270., 13152928.]],

            [[ 8883886.,  9884376.],
              [12114390., 13179168.]]], dtype=float32)
```

1.1.5 Benchmark

It clearly shows the summation done with the basic algorithm is the slowest.

```
[15]: from mlproduct.onnxrt.validate.validate_helper import measure_time
from tqdm import tqdm
from pandas import DataFrame

def raw_product(m1, m2, m3):
    N = m1.shape[0]
    res = numpy.zeros((N, N, N))
    for a in range(0, N):
        for b in range(0, N):
            for c in range(0, N):
                for d in range(0, N):
                    for e in range(0, N):
                        for f in range(0, N):
                            res[e, b, c] += m1[b, a, c] * m2[c, d] * m3[d, e, f]
    return res

def benchmark0(equation):
    sess = None
    sess2 = None
    seq = None
    seq2 = None

    results = []
    for N in tqdm([2, 3, 4, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60]):
        m1 = numpy.random.randn(N, N, N)
        m2 = numpy.random.randn(N, N)
        m3 = numpy.random.randn(N, N, N)

        if seq is None:
            seq = decompose_einsum_equation(equation, clean=True)
        if seq2 is None:
            seq2 = decompose_einsum_equation(equation, clean=True, strategy='numpy')
        if sess is None:
            model = make_model1(equation)
            sess = InferenceSession(model.SerializeToString())
        if sess2 is None:
            onx = seq2.to_onnx("Y", "X1", "X2", "X3", dtype=numpy.float32)
            sess2 = InferenceSession(onx.SerializeToString())

        res = measure_time(lambda x: numpy.einsum(equation, *x, optimize=True),
                           [m1, m2, m3],
                           repeat=10, number=10)
```

```

res['name'] = "numpy.einsum"
res["N"] = N
results.append(res)

if N <= 4:
    res = measure_time(lambda x: raw_product(*x),
                        [m1, m2, m3],
                        repeat=10, number=10)
    res['name'] = "raw_product"
    res["N"] = N
    results.append(res)

res = measure_time(lambda x: apply_einsum_sequence(seq, *x),
                    [m1, m2, m3],
                    repeat=10, number=10)

res['name'] = "custom_einsum"
res["N"] = N
results.append(res)

res = measure_time(lambda x: apply_einsum_sequence(seq, *x, matmul_impl="pyf"),
                    [m1, m2, m3],
                    repeat=10, number=10)
res['name'] = "dec-matmul"
res["N"] = N
results.append(res)

res = measure_time(lambda x: apply_einsum_sequence(seq2, *x,
matmul_impl="pyf"),
                    [m1, m2, m3],
                    repeat=10, number=10)
res['name'] = "dec-batch_dot"
res["N"] = N
results.append(res)

res = measure_time(lambda x: sess.run(None, {'X': x[0], 'Y': x[1], 'Z': x[2]}),
                    [m1.astype(numpy.float32), m2.astype(numpy.float32),
                     m3.astype(numpy.float32)],
                    repeat=10, number=10)
res['name'] = "ort-einsum"
res["N"] = N
results.append(res)

res = measure_time(lambda x: sess2.run(None, {'X1': x[0], 'X2': x[1], 'X3':
x[2]}),
                    [m1.astype(numpy.float32), m2.astype(numpy.float32),
                     m3.astype(numpy.float32)],
                    repeat=10, number=10)
res['name'] = "ort-matmul"
res["N"] = N
results.append(res)
return DataFrame(results)

```

```
df = benchmark0("bac,cd,def->ebc")
df.tail()
```

C:\xavierdupre_home_github_fork\scikit-learn\sklearn\experimental\enable_hist_gradient_boosting.py:16: UserWarning: Since version 1.0, it is not needed to import enable_hist_gradient_boosting anymore. HistGradientBoostingClassifier and HistGradientBoostingRegressor are now stable and can be normally imported from sklearn.ensemble.

```
warnings.warn(
100%|_ _ _ _ _ _ _ _ _ _ | 14/14 [00:20<00:00, 1.47s/it]
```


```
[15]:      average  deviation  min_exec  max_exec  repeat  number  total  \
82  0.065132   0.001338  0.063801  0.068927     10     10  0.651318
83  0.051615   0.001206  0.049987  0.053465     10     10  0.516154
84  0.062689   0.003658  0.058949  0.073073     10     10  0.626888
85  0.009917   0.000274  0.009737  0.010686     10     10  0.099166
86  0.015518   0.001107  0.014413  0.018179     10     10  0.155178
```

```
      name  N
82  custom_einsum  60
83      dec-matmul  60
84  dec-batch_dot  60
85      ort-einsum  60
86      ort-matmul  60
```

```
[16]: import matplotlib.pyplot as plt

piv = df.pivot("N", "name", "average")
piv2 = piv.copy()
np = piv["numpy.einsum"]
for c in piv2.columns:
    piv2[c] /= np

fig, ax = plt.subplots(1, 2, figsize=(14, 6))
piv.plot(logy=True, logx=True, ax=ax[0])
ax[0].set_title("Benchmark einsum function\nbac,cd,def->ebc")
piv2.plot(logy=True, logx=True, ax=ax[1])
ax[1].set_title("Benchmark einsum function\n(ratio, baseline=numpy)");
```



Version `dec-matmul` is an implementation based on the decomposition of a simplified einsum into a sequence of transpose, reshape, (batch_)dot or mul operations. This decomposition is converted into ONNX and executed with `onnxruntime`, version `ort-matmul`. Both versions are faster than the numpy optimized version.

1.2 Another example with `bsnh,btnh=bnts`

Another case, more frequent in deep learning.

1.2.1 Decomposition of `bsnh,btnh=bnts`

```
[17]: seq2 = decompose_einsum_equation("bsnh,btnh->bnts", strategy='numpy', clean=True)
      RenderJsDot(seq2.to_dot(size=7))
```

```
[17]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x1c6d8eae100>
```

1.2.2 ONNX version

```
[18]: onx2 = seq2.to_onnx("Y", "X1", "X2", dtype=numpy.float32)
      %onnxview onx2
```

```
[18]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x1c6d91382b0>
```

1.2.3 Benchmark

```
[19]: def make_model2(equation):
      model = helper.make_model(
          opset_imports=[helper.make_operatorsetid('', 13)],
          graph=helper.make_graph(
              name='einsum_test',
              inputs=[helper.make_tensor_value_info("X", onnx.TensorProto.FLOAT, None),
                      helper.make_tensor_value_info("Y", onnx.TensorProto.FLOAT, None)],
              outputs=[helper.make_tensor_value_info("A", onnx.TensorProto.FLOAT, None)],
              nodes=
```

```

        helper.make_node("Einsum", ["X", "Y"], ["A"], equation=equation)
    ]
)
)
return model

def benchmark(equation, second_input_size=4):
    sess = None
    sess2 = None
    seq = None
    seq2 = None

    results = []
    for N in tqdm([2, 3, 4, 10, 20, 30, 40]):
        m1 = numpy.random.randn(10, N, N, N)
        m2 = numpy.random.randn(10 * N ** (second_input_size-1)).reshape((10, ) + (N,
↪) * (second_input_size-1))

        if seq is None:
            seq = decompose_einsum_equation(equation, clean=True)
        if seq2 is None:
            seq2 = decompose_einsum_equation(equation, clean=True, strategy='numpy')
        if sess is None:
            model = make_model2(equation)
            sess = InferenceSession(model.SerializeToString())
        if sess2 is None:
            onx = seq2.to_onnx("Y", "X1", "X2", dtype=numpy.float32)
            sess2 = InferenceSession(onx.SerializeToString())

        res = measure_time(lambda x: numpy.einsum(equation, *x, optimize=True),
                           [m1, m2],
                           repeat=10, number=10)

        res['name'] = "numpy.einsum"
        res["N"] = N
        results.append(res)

        res = measure_time(lambda x: apply_einsum_sequence(seq, *x),
                           [m1, m2],
                           repeat=10, number=10)
        res['name'] = "custom_einsum"
        res["N"] = N
        results.append(res)

        res = measure_time(lambda x: apply_einsum_sequence(seq, *x, matmul_impl="pyf"),
                           [m1, m2],
                           repeat=10, number=10)
        res['name'] = "dec-matmul"
        res["N"] = N
        results.append(res)

```



```

        res = measure_time(lambda x: apply_einsum_sequence(seq2, *x,
matmul_impl="pyf"),
                           [m1, m2],
                           repeat=10, number=10)
    res['name'] = "dec-batch_dot"
    res["N"] = N
    results.append(res)

    res = measure_time(lambda x: sess.run(None, {'X': x[0], 'Y': x[1]}),
                       [m1.astype(numpy.float32), m2.astype(numpy.float32),
                        m3.astype(numpy.float32)],
                       repeat=10, number=10)
    res['name'] = "ort-einsum"
    res["N"] = N
    results.append(res)

    res = measure_time(lambda x: sess2.run(None, {'X1': x[0], 'X2': x[1]}),
                       [m1.astype(numpy.float32), m2.astype(numpy.float32),
                        m3.astype(numpy.float32)],
                       repeat=10, number=10)
    res['name'] = "ort-matmul"
    res["N"] = N
    results.append(res)
    return DataFrame(results)

df = benchmark("bsnh,btnh->bnts")
df.tail()

```

100%| ██████████ | 7/7 [00:13<00:00, 1.93s/it]

```

[19]:
   average  deviation  min_exec  max_exec  repeat  number  total \
37  0.229418   0.020792  0.217997  0.291032      10      10  2.294175
38  0.160575   0.005435  0.150772  0.167411      10      10  1.605746
39  0.112844   0.011305  0.102173  0.141890      10      10  1.128436
40  0.051181   0.003533  0.047244  0.057054      10      10  0.511815
41  0.078827   0.008735  0.067893  0.099156      10      10  0.788271

```

```

      name  N
37  custom_einsum  40
38    dec-matmul  40
39  dec-batch_dot  40
40    ort-einsum  40
41    ort-matmul  40

```


```

[20]: piv = df.pivot("N", "name", "average")
      piv2 = piv.copy()
      np = piv["numpy.einsum"]
      for c in piv2.columns:
          piv2[c] /= np

      fig, ax = plt.subplots(1, 2, figsize=(14, 6))

```

```
piv.plot(logy=True, logx=True, ax=ax[0])
ax[0].set_title("Benchmark einsum function\nbsnh,btnh->bnts")
piv2.plot(logy=True, logx=True, ax=ax[1])
ax[1].set_title("Benchmark einsum function\nn(ratio, baseline=numpy)");
```



1.2.4 Permutation

Einsum's algorithm started by aligning all matrices involved in the computation to the same dimension in the same order. But which order is the best, that's the question.

```
[21]: equation = "bsnh,btnh->bnts"
letters = list(sorted(set([c for c in equation if "a" <= c < "z"])))
letters
```

```
[21]: ['b', 'h', 'n', 's', 't']
```

```
[22]: from itertools import permutations

def benchmark_perm(equation, number=5, second_input_size=4, repeat=3, N=15):

    def n_operator(seq, name):
        n = 0
        for op in seq:
            if op.name == name:
                n += 1
        return n

    def n_onnx_op(onx, name):
        n = 0
        for op in onx.graph.node:
            if op.op_type == name:
```

```

        n += 1
    return n

def get_kind(seq):
    n = 0
    for op in seq:
        if op.name == 'batch_dot':
            return op.get_dot_kind()
    return None

m1 = numpy.random.randn(N, N, N, N)
m2 = numpy.random.randn(N ** second_input_size).reshape((N, ) * second_input_size)

results = []
for perm in tqdm(list(permutations(letters))):
    replace = {d: c for c, d in zip(letters, perm)}
    eq = equation
    for k, v in replace.items():
        eq = eq.replace(k, v.upper())
    eq = eq.lower()

    seq = decompose_einsum_equation(eq, clean=True)
    seq2 = decompose_einsum_equation(eq, clean=True, strategy='numpy')
    model = make_model2(eq)
    sess = InferenceSession(model.SerializeToString())
    onx = seq2.to_onnx("Y", "X1", "X2", dtype=numpy.float32)
    sess2 = InferenceSession(onx.SerializeToString())

    n_tra = n_operator(seq2, 'transpose')
    n_tra_onnx = n_onnx_op(onx, 'Transpose')
    n_gemm_onnx = n_onnx_op(onx, 'Gemm')
    kind = get_kind(seq2)

    res = measure_time(lambda x: numpy.einsum(eq, *x, optimize=True),
                        [m1, m2],
                        repeat=repeat, number=number)

    res['name'] = "numpy.einsum"
    res["N"] = N
    res["eq"] = eq
    results.append(res)

    res = measure_time(lambda x: apply_einsum_sequence(seq, *x),
                        [m1, m2],
                        repeat=repeat, number=number)

    res['name'] = "custom_einsum"
    res["N"] = N
    res["eq"] = eq
    res['transpose'] = n_tra
    res['kind'] = kind
    results.append(res)

```

```

    res = measure_time(lambda x: apply_einsum_sequence(seq, *x, matmul_impl="pyf"),
                        [m1, m2],
                        repeat=repeat, number=number)
    res['name'] = "dec-matmul"
    res["N"] = N
    res["eq"] = eq
    res['transpose'] = n_tra
    res['kind'] = kind
    results.append(res)

    res = measure_time(lambda x: apply_einsum_sequence(seq2, *x,
matmul_impl="pyf"),
                        [m1, m2],
                        repeat=repeat, number=number)
    res['name'] = "dec-batch_dot"
    res["N"] = N
    res["eq"] = eq
    res['transpose'] = n_tra
    res['kind'] = kind
    results.append(res)

    res = measure_time(lambda x: sess.run(None, {'X': x[0], 'Y': x[1]}),
                        [m1.astype(numpy.float32), m2.astype(numpy.float32),
                        m3.astype(numpy.float32)],
                        repeat=repeat, number=number)
    res['name'] = "ort-einsum"
    res["N"] = N
    res["eq"] = eq
    res['transpose'] = n_tra_onnx
    res['gemm'] = n_gemm_onnx
    results.append(res)

    res = measure_time(lambda x: sess2.run(None, {'X1': x[0], 'X2': x[1]}),
                        [m1.astype(numpy.float32), m2.astype(numpy.float32),
                        m3.astype(numpy.float32)],
                        repeat=repeat, number=number)
    res['name'] = "ort-matmul"
    res["N"] = N
    res["eq"] = eq
    res['transpose'] = n_tra_onnx
    res['gemm'] = n_gemm_onnx
    results.append(res)
return DataFrame(results)

df = benchmark_perm("bsnh,btnh->bnts", number=4)
df.tail()

```

100%| ██████████ | 120/120 [00:11<00:00, 10.23it/s]

```

[22]:      average  deviation  min_exec  max_exec  repeat  number  total  \
715  0.006162   0.000038  0.006128  0.006216         3         4  0.018485

```

716	0.002343	0.000046	0.002294	0.002405	3	4	0.007029
717	0.001645	0.000035	0.001610	0.001694	3	4	0.004934
718	0.000833	0.000015	0.000820	0.000853	3	4	0.002498
719	0.001251	0.000012	0.001238	0.001268	3	4	0.003753

	name	N	eq	transpose	kind	gemm
715	custom_einsum	15	thns,tbns->tnbh	3.0	NN	NaN
716	dec-matmul	15	thns,tbns->tnbh	3.0	NN	NaN
717	dec-batch_dot	15	thns,tbns->tnbh	3.0	NN	NaN
718	ort-einsum	15	thns,tbns->tnbh	4.0	NaN	0.0
719	ort-matmul	15	thns,tbns->tnbh	4.0	NaN	0.0

```
[23]: df = df.sort_values("average").reset_index(drop=True)
df.head()
```

```
[23]:      average  deviation  min_exec  max_exec  repeat  number  total  \
0  0.000758   0.000015   0.000738   0.000771        3        4  0.002275
1  0.000770   0.000023   0.000739   0.000793        3        4  0.002310
2  0.000778   0.000020   0.000758   0.000806        3        4  0.002334
3  0.000783   0.000021   0.000760   0.000812        3        4  0.002350
4  0.000784   0.000011   0.000774   0.000799        3        4  0.002351
```

	name	N	eq	transpose	kind	gemm
0	ort-matmul	15	hsnt,hbnt->hnbs	4.0	NaN	0.0
1	ort-matmul	15	hnts,hbts->htbn	4.0	NaN	0.0
2	ort-matmul	15	bnst,bhst->bshn	4.0	NaN	0.0
3	ort-matmul	15	bnht,bsht->bhsn	4.0	NaN	0.0
4	ort-matmul	15	hnst,hbst->hsbn	4.0	NaN	0.0

```
[24]: df.tail()
```

```
[24]:      average  deviation  min_exec  max_exec  repeat  number  total  \
715  0.011529   0.000882   0.010456   0.012617        3        4  0.034587
716  0.011548   0.000422   0.010967   0.011953        3        4  0.034644
717  0.013971   0.001984   0.012279   0.016754        3        4  0.041912
718  0.014765   0.001483   0.013366   0.016818        3        4  0.044295
719  0.015813   0.002921   0.012546   0.019636        3        4  0.047438
```

	name	N	eq	transpose	kind	gemm
715	custom_einsum	15	sbnt,shnt->snhb	3.0	NN	NaN
716	custom_einsum	15	htsb,hnsb->hsnt	3.0	NN	NaN
717	custom_einsum	15	nbsh,ntsh->nstb	3.0	NN	NaN
718	numpy.einsum	15	bnsh,btsh->bstn	NaN	NaN	NaN
719	numpy.einsum	15	nbsh,ntsh->nstb	NaN	NaN	NaN

```
[25]: piv = df.pivot("eq", "name", "average").sort_values("numpy.einsum")

fig, ax = plt.subplots(1, 1, figsize=(14, 6))
piv.plot(logy=True, logx=True, ax=ax)
ax.set_title("Benchmark einsum function - bsnh,btnh->bnts");
```




```

38      dec-matmul    40
39  dec-batch_dot    40
40      ort-einsum    40
41      ort-matmul    40

```

```
[30]: piv = df.pivot("N", "name", "average")
      piv2 = piv.copy()
      np = piv["numpy.einsum"]
      for c in piv2.columns:
          piv2[c] /= np

      fig, ax = plt.subplots(1, 2, figsize=(14, 6))
      piv.plot(logy=True, logx=True, ax=ax[0])
      ax[0].set_title("Benchmark einsum function\nbsnh,ctnh->nts")
      piv2.plot(logy=True, logx=True, ax=ax[1])
      ax[1].set_title("Benchmark einsum function\nn(ratio, baseline=numpy)");
```



1.3.2 Benchmark permutation

```
[31]: df = benchmark_perm("bsnh,ctnh->nts", number=2, repeat=3, N=10)
```

```
100%| . . . . . | 120/120 [00:06<00:00, 17.41it/s]
```

```
[32]: df = df.sort_values("average").reset_index(drop=True)
df.head()
```

```
[32]:
```

	average	deviation	min_exec	max_exec	repeat	number	total	\
0	0.000125	0.000008	0.000118	0.000136	3	2	0.000374	
1	0.000126	0.000007	0.000119	0.000136	3	2	0.000377	
2	0.000141	0.000006	0.000136	0.000150	3	2	0.000422	
3	0.000141	0.000007	0.000135	0.000151	3	2	0.000423	

```
4 0.000144 0.000007 0.000138 0.000154 3 2 0.000432
```


	name	N	eq	transpose	kind	gemm
0	ort-matmul	10	bnst,chst->shn	4.0	NaN	0.0
1	ort-matmul	10	bhst,cnst->snh	4.0	NaN	0.0
2	ort-matmul	10	hbst,cnst->snb	5.0	NaN	0.0
3	ort-matmul	10	nbst,chst->shb	5.0	NaN	0.0
4	ort-matmul	10	btns,chns->nht	5.0	NaN	0.0

```
[33]: set(df['transpose'].dropna()), set(df['gemm'].dropna()), set(df['kind'].dropna())
```

```
[33]: ({3.0, 4.0, 5.0, 6.0}, {0.0}, {'NN'})
```

```
[34]: piv = df.pivot("eq", "name", "average").sort_values("numpy.einsum")
```

```
fig, ax = plt.subplots(1, 1, figsize=(14, 6))
piv.plot(logy=True, logx=True, ax=ax)
ax.set_title("Benchmark einsum function");
```



1.3.3 Best permutation

One of the best permutation is `bnst,chst->shn`.

```
[35]: seq4 = decompose_einsum_equation("bnst,chst->shn", strategy='numpy', clean=True)
RenderJsDot(seq4.to_dot(size=7))
```

```
[35]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x1c6d8e73640>
```

```
[36]: onx4 = seq4.to_onnx("Y", "X1", "X2", dtype=numpy.float32)
%onnxview onx4
```

```
[36]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x1c6d9428d90>
```


[37] :