

# onnx\_pdist

March 10, 2022

## 1 Pairwise distances with ONNX (pdist)

Function `pdist` computes pairwise distances between observations in  $n$ -dimensional space. It is not that difficult to convert that into *ONNX* when the dimension of the input is always the same. What if not?

```
[1]: from jyquickhelper import add_notebook_menu
      add_notebook_menu()
```

```
[1]: <IPython.core.display.HTML object>
```

```
[2]: %load_ext mlproduct
```

The `mlproduct` extension is already loaded. To reload it, use:  
`%reload_ext mlproduct`

### 1.1 Function `pdist`

The function `pdist` distances. Let's denote a list of vectors  $(X_1, \dots, X_n)$ , function `pdist` returns the matrix  $D = (d_{ij})$  where  $d_{ij} = \text{dist}(X_i, X_j) = \|X_i - X_j\|^2$ .

```
[3]: import numpy
      from scipy.spatial.distance import pdist, squareform

      M = numpy.array([[0, 1],
                      [1, 2],
                      [0.1, 1.1],
                      [2, 2]], dtype=float)

      d1 = squareform(pdist(M, metric='sqeuclidean'))
      d1
```

```
[3]: array([[0. , 2. , 0.02, 5. ],
          [2. , 0. , 1.62, 1. ],
          [0.02, 1.62, 0. , 4.42],
          [5. , 1. , 4.42, 0. ]])
```

The two following functions are implemented to reduce the number of allocations the algorithm requires.

```
[4]: def custom_pdist(M):
      n = M.shape[0]
      res = numpy.zeros((n, n))
      buffer = numpy.empty(M.shape)
      for i in range(n):
```

```

        numpy.subtract(M, M[i], out=buffer) # broadcasted subtraction
        numpy.square(buffer, out=buffer)
        res[i, :] = numpy.sum(buffer, axis=1)
    return res

d2 = custom_pdist(M)
d2

```

```

[4]: array([[0. , 2. , 0.02, 5. ],
           [2. , 0. , 1.62, 1. ],
           [0.02, 1.62, 0. , 4.42],
           [5. , 1. , 4.42, 0. ]])

```

This function computes  $n^2$  distances whereas only  $\frac{n(n-1)}{2}$  are necessary since the final matrix is symmetric. Let's change the implementation to reflect that.

```

[5]: def custom_pdist_lower(M):
    n = M.shape[0]
    res = numpy.zeros((n, n))
    buffer = numpy.empty((M.shape[0]-1, M.shape[1]))
    a = numpy.empty(M.shape[0])
    for i in range(1, n):
        numpy.subtract(M[:i], M[i], out=buffer[:i]) # broadcasted subtraction
        numpy.square(buffer[:i], out=buffer[:i])
        numpy.sum(buffer[:i], axis=1, out=a[:i])
        res[:i, i] = a[:i]
        res[i, :i] = a[:i]
    return res

d3 = custom_pdist_lower(M)
d3

```

```

[5]: array([[0. , 2. , 0.02, 5. ],
           [2. , 0. , 1.62, 1. ],
           [0.02, 1.62, 0. , 4.42],
           [5. , 1. , 4.42, 0. ]])

```

## 1.2 Loop mechanism in ONNX

Operator [Loop](#) seems appropriate but it is just a loop whereas [Scan](#) holds accumulator. The first graph is what is repeated inside the loop.

```

[6]: from skl2onnx.algebra.onnx_ops import OnnxAdd, OnnxIdentity, OnnxScan
    from skl2onnx.common.data_types import FloatTensorType

    initial = numpy.array([0, 0]).astype(numpy.float32).reshape((2,))
    x = numpy.array([1, 2, 3, 4, 5, 6]).astype(numpy.float32).reshape((3, 2))

    add_node = OnnxAdd('sum_in', 'next', output_names=['sum_out'], op_version=12)
    id_node = OnnxIdentity(add_node, output_names=['scan_out'], op_version=12)

    scan_body = id_node.to_onnx(
        {'sum_in': initial, 'next': initial},
        outputs=[('sum_out', FloatTensorType())],
    )

```

```

        ('scan_out', FloatTensorType()))

# add -l 1 if nothing shows up
%onnxview scan_body

```

[6]: <jyquickhelper.jspy.render\_nb\_js\_dot.RenderJsDot at 0x234da711a90>

The operator `Scan` repeats this graph a couple of times. `sum_in` is an accumulator, `next` is the iterated row from the input matrix.

```

[7]: node = OnnxScan('initial', 'x', output_names=['y', 'z'],
                    num_scan_inputs=1, body=scan_body.graph)

model_def = node.to_onnx(
    {'initial': initial, 'x': x},
    outputs=[('y', FloatTensorType()),
             ('z', FloatTensorType())])

# add -l 1 if nothing shows up
%onnxview model_def

```

[7]: <jyquickhelper.jspy.render\_nb\_js\_dot.RenderJsDot at 0x234da3810a0>

All together in the same graph.

```

[8]: # add -l 1 if nothing shows up
%onnxview model_def -r 1

```

[8]: <jyquickhelper.jspy.render\_nb\_js\_dot.RenderJsDot at 0x234dc1e7880>

```

[9]: from mlprodict.onnxrt import OnnxInference
oinf = OnnxInference(model_def)
res = oinf.run({'initial': initial, 'x': x})
res['y']

```

[9]: array([ 9., 12.], dtype=float32)

```

[10]: res['z']

```

```

[10]: array([[ 1.,  2.],
            [ 4.,  6.],
            [ 9., 12.]], dtype=float32)

```

### 1.3 Back to pdist

`sklearn-onnx` implements function `pdist` with `ONNX` operators. The parameter `inputs=[('x', FloatTensorType())]` tells the method `to_onnx` that the dimension of the inputs is not fixed and should not be checked.

```

[11]: # from skl2onnx.algebra.complex_functions import squareform_pdist_

from collections import OrderedDict
from skl2onnx.algebra.onnx_ops import (
    OnnxSub, OnnxReduceSumSquare, OnnxSqueeze,
    OnnxIdentity, OnnxScan)

```

```

from skl2onnx.common.data_types import FloatTensorType
from mlproduct.tools import get_opset_number_from_onnx

def squareform_pdist(X, **kwargs):
    """Returns the ONNX graph which computes
    ``squareform(pdist(X, metric='sqeuclidean'))``."""

    # The subgraph executed at every iteration.
    opv = get_opset_number_from_onnx()
    diff = OnnxSub('next_in', 'next', output_names=['diff'], op_version=opv)
    id_next = OnnxIdentity('next_in', output_names=['next_out'], op_version=opv)
    norm = OnnxReduceSumSquare(diff, output_names=['norm'], axes=[1], op_version=opv)
    flat = OnnxSqueeze(norm, numpy.array([1], dtype=numpy.int64),
                       output_names=['scan_out'], op_version=opv)
    scan_body = id_next.to_onnx(
        OrderedDict([('next_in', FloatTensorType()),
                     ('next', FloatTensorType())]),
        # Size must be empty otherwise onnxruntime fails
        # at execution time if it receives a matrix
        # with a different shape. With 'None', the same ONNX graph
        # can compute pairwise distance for any shape.
        outputs=[('next_out', FloatTensorType([None, None])),
                 ('scan_out', FloatTensorType([None]))],
        other_outputs=[flat])

    # The loop.
    # 'scan0_{idself}' means the variable name will include
    # id(OnnxScan), this is needed if squareform_pdist is used
    # twice in the same graph.
    node = OnnxScan(X, X, output_names=['scan0_{idself}', 'scan1_{idself}'],
                    num_scan_inputs=1, body=scan_body.graph, op_version=opv,
                    **kwargs)

    return node[1]

opv = get_opset_number_from_onnx()
onnx_fct = OnnxIdentity(squareform_pdist('x'), output_names='Y', op_version=opv)
model_def = onnx_fct.to_onnx(inputs=[('x', FloatTensorType())])

# add -l 1 if nothing shows up
%onnxview model_def

```

[11]: <jyquickhelper.jspy.render\_nb\_js\_dot.RenderJsDot at 0x234da7785e0>

```

[12]: from collections import OrderedDict
from skl2onnx.algebra.onnx_ops import (
    OnnxSub, OnnxReduceSumSquare, OnnxSqueeze,
    OnnxIdentity, OnnxScan)
from skl2onnx.common.data_types import FloatTensorType
from mlproduct.tools import get_opset_number_from_onnx

def squareform_pdist(X, **kwargs):

```

```

# The subgraph executed at every iteration.
opv = get_opset_number_from_onnx()
diff = OnnxSub('next_in', 'next', output_names=['diff'], op_version=opv)
id_next = OnnxIdentity('next_in', output_names=['next_out'], op_version=opv)
norm = OnnxReduceSumSquare(diff, output_names=['norm'], axes=[1], op_version=opv)
flat = OnnxSqueeze(norm, numpy.array([1], dtype=numpy.int64),
                    output_names=['scan_out'], op_version=opv)
scan_body = id_next.to_onnx(
    OrderedDict([('next_in', FloatTensorType()),
                  ('next', FloatTensorType())]),
    outputs=[('next_out', FloatTensorType([None, None])),
              ('scan_out', FloatTensorType([None]))],
    other_outputs=[flat])

# The loop.
node = OnnxScan(X, X, output_names=['scan0_{idself}', 'scan1_{idself}'],
                 num_scan_inputs=1, body=scan_body.graph, op_version=opv,
                 **kwargs)

return node[1]

opv = get_opset_number_from_onnx()
onnx_fct = OnnxIdentity(squareform_pdist('x'), output_names='Y', op_version=opv)
model_def = onnx_fct.to_onnx(inputs=[('x', FloatTensorType())])

```

Notice the double arrow. Input  $x$  is used twice, once as an permanent state involved in broadcasted substract, another time to iterator rows. On the other side, the first output of operator *Scan* is a permanent state equal to the input, the second one is an aggregation of results produced at each iteration. Each of those produces a row of a final matrix.

```

[13]: oinf = OnnxInference(model_def)
      body = oinf['Sc_Scan', 'body']

# add -l 1 if nothing shows up
      %onnxview body.g

```

```
[13]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x234dc3b2a00>
```

All together.

```

[14]: # add -l 1 if nothing shows up
      %onnxview model_def -r 1

```

```
[14]: <jyquickhelper.jspy.render_nb_js_dot.RenderJsDot at 0x234dc3b23d0>
```

Let's now execute the graph and compare it with the original graph.

```

[15]: d1 = squareform(pdist(M, metric='sqeuclidean'))
      d1

```

```

[15]: array([[0.   , 2.   , 0.02, 5.   ],
            [2.   , 0.   , 1.62, 1.   ],
            [0.02, 1.62, 0.   , 4.42],
            [5.   , 1.   , 4.42, 0.   ]])

```

```

[16]: oinf.run({'x': M})['Y']

```

```
[16]: array([[0. , 2. , 0.02, 5. ],
           [2. , 0. , 1.62, 1. ],
           [0.02, 1.62, 0. , 4.42],
           [5. , 1. , 4.42, 0. ]])
```

```
[17]: %timeit squareform(pdist(M, metric='sqeuclidean'))
```

9.31  $\mu$ s  $\pm$  423 ns per loop (mean  $\pm$  std. dev. of 7 runs, 100,000 loops each)

```
[18]: %timeit custom_pdist(M)
```

35.1  $\mu$ s  $\pm$  1.52  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 10,000 loops each)

```
[19]: %timeit custom_pdist_lower(M)
```

34.2  $\mu$ s  $\pm$  2.18  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 10,000 loops each)

```
[20]: %timeit oinf.run({'x': M})['Y']
```

177  $\mu$ s  $\pm$  11.3  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 10,000 loops each)

```
[21]: M32 = M.astype(numpy.float32)
```

```
[22]: from mlprodict.tools import get_ir_version_from_onnx
model_def.ir_version = get_ir_version_from_onnx()
```

```
[23]: oinftrt = OnnxInference(model_def, runtime="onnxruntime1")
oinftrt.run({'x': M32})['Y']
```

No CUDA runtime is found, using CUDA\_HOME='C:\Program Files\NVIDIA GPU Computing Toolkit\CUDA\v11.5'

```
[23]: array([[0. , 2. , 0.02000001, 5. ,
           [2. , 0. , 1.6199999 , 1. ,
           [0.02000001, 1.6199999 , 0. , 4.42 ,
           [5. , 1. , 4.42 , 0. ,
           dtype=float32])
```

```
[24]: %timeit oinftrt.run({'x': M32})['Y']
```

43.1  $\mu$ s  $\pm$  4.32  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 10,000 loops each)

## 1.4 Benchmark

```
[25]: from timeit import Timer
```

```
def measure_time(name, stmt, context, repeat=10, number=10):
    tim = Timer(stmt, globals=context)
```

```

res = numpy.array(tim.repeat(repeat=repeat, number=number))
res /= number
mean = numpy.mean(res)
dev = numpy.mean(res ** 2)
dev = (dev - mean**2) ** 0.5
return dict(average=mean, deviation=dev, min_exec=numpy.min(res),
            max_exec=numpy.max(res), repeat=repeat, number=number,
            nrows=context['M'].shape[0], ncols=context['M'].shape[1],
            name=name)

measure_time("scipy", "squareform(pdist(M, metric='sqeuclidean'))",
            context={'squareform': squareform, 'M': M,
                    'pdist': pdist})

```

```

[25]: {'average': 4.233300000009876e-05,
      'deviation': 2.7235873787981297e-05,
      'min_exec': 1.8629999999575375e-05,
      'max_exec': 0.00010153999999999997,
      'repeat': 10,
      'number': 10,
      'nrows': 4,
      'ncols': 2,
      'name': 'scipy'}

```

```

[26]: from tqdm import trange

def generator():
    for feat in [5, 10, 50, 100]:
        for n in [5, 10, 20, 50, 100, 400, 1000]:
            if n <= 500 or feat <= 10:
                yield feat, n

all_values = list(generator())

rows = []

with trange(len(all_values)) as t:
    for i in t:
        feat, n = all_values[i]
        t.set_description("feat=%d n=%d" % (feat, n))
        M = numpy.random.rand(n, feat)

        context = {'squareform': squareform, 'M': M, 'pdist': pdist}
        res = measure_time("scipy", "squareform(pdist(M, metric='sqeuclidean'))",
            context=context)
        res['dimres'] = squareform(pdist(M, metric='sqeuclidean')).shape[0]
        rows.append(res)

        context = {'M': M, 'custom_pdist': custom_pdist}
        res = measure_time("numpy", "custom_pdist(M)", context=context)
        res['dimres'] = custom_pdist(M).shape[0]
        rows.append(res)

```

```

context = {'M': M, 'custom_pdist_lower': custom_pdist_lower}
res = measure_time("numpy-lower", "custom_pdist_lower(M)", context=context)
res['dimres'] = custom_pdist_lower(M).shape[0]
rows.append(res)

context = {'oinf': oinf, 'M': M}
res = measure_time("onnx-py", "oinf.run({'x': M})['Y']", context=context)
res['dimres'] = oinf.run({'x': M})['Y'].shape[0]
rows.append(res)

M32 = M.astype(numpy.float32)
context = {'oinfrt': oinfrt, 'M': M32}
res = measure_time("onnx-rt", "oinfrt.run({'x': M})['Y']", context=context)
res['dimres'] = oinfrt.run({'x': M32})['Y'].shape[0]
rows.append(res)

from pandas import DataFrame
df = DataFrame(rows)
df.head()

```

feat=100 n=400: 100%| 26/26 [01:20<00:00, 3.10s/it]

```

[26]:
  average  deviation  min_exec  max_exec  repeat  number  nrows  ncols  \
0  0.000015  0.000005  0.000010  0.000025      10      10      5      5
1  0.000106  0.000023  0.000065  0.000138      10      10      5      5
2  0.000053  0.000005  0.000048  0.000064      10      10      5      5
3  0.000240  0.000017  0.000219  0.000273      10      10      5      5
4  0.000053  0.000008  0.000046  0.000072      10      10      5      5

   name  dimres
0  scipy      5
1  numpy      5
2  numpy-lower  5
3  onnx-py      5
4  onnx-rt      5

```

```

[27]: from pandas import pivot_table
piv = pivot_table(df, index=["nrows"], columns= ['ncols', 'name'], values='average')
piv.head().T

```

```

[27]:
nrows      5      10      20      50      100
ncols name
5  numpy      0.000106  0.000108  0.000193  0.000464  0.001121
   numpy-lower  0.000053  0.000099  0.000225  0.000520  0.001190
   onnx-py      0.000240  0.000407  0.000797  0.002581  0.003790
   onnx-rt      0.000053  0.000071  0.000118  0.000306  0.000766
   scipy      0.000015  0.000011  0.000014  0.000020  0.000044
10  numpy      0.000067  0.000094  0.000194  0.000569  0.001441
   numpy-lower  0.000044  0.000093  0.000189  0.000591  0.001209
   onnx-py      0.000226  0.000379  0.000751  0.001945  0.004731
   onnx-rt      0.000048  0.000072  0.000144  0.000329  0.000995
   scipy      0.000013  0.000013  0.000016  0.000023  0.000071

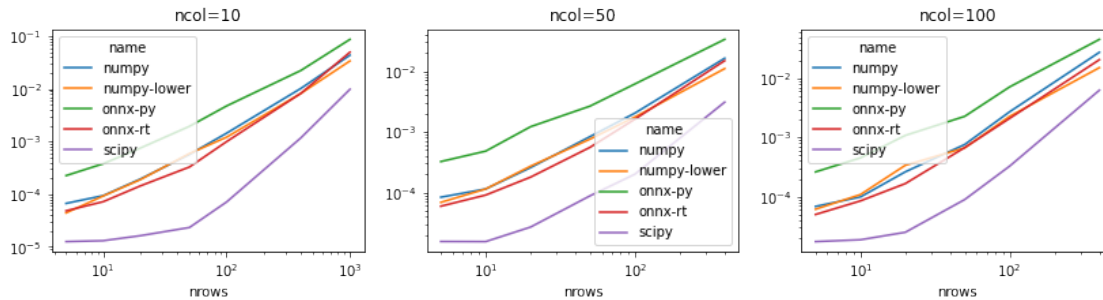
```



50	numpy	0.000084	0.000114	0.000257	0.000833	0.002031
	numpy-lower	0.000069	0.000114	0.000272	0.000757	0.001749
	onnx-py	0.000323	0.000480	0.001214	0.002648	0.006138
	onnx-rt	0.000059	0.000091	0.000179	0.000554	0.001614
	scipy	0.000016	0.000016	0.000027	0.000088	0.000200
100	numpy	0.000068	0.000098	0.000262	0.000759	0.002712
	numpy-lower	0.000061	0.000108	0.000338	0.000666	0.002270
	onnx-py	0.000261	0.000451	0.001082	0.002272	0.007142
	onnx-rt	0.000050	0.000084	0.000166	0.000672	0.002097
	scipy	0.000017	0.000019	0.000025	0.000089	0.000327

```
[28]: %matplotlib inline
```

```
[29]: import matplotlib.pyplot as plt
fig, ax = plt.subplots(1, 3, figsize=(14, 3))
for i, ncol in enumerate([10, 50, 100]):
    piv = df[df.ncols==ncol].pivot("nrows", "name", "average")
    piv.plot(ax=ax[i], logy=True, logx=True)
    ax[i].set_title("ncol=%d" % ncol)
ax;
```

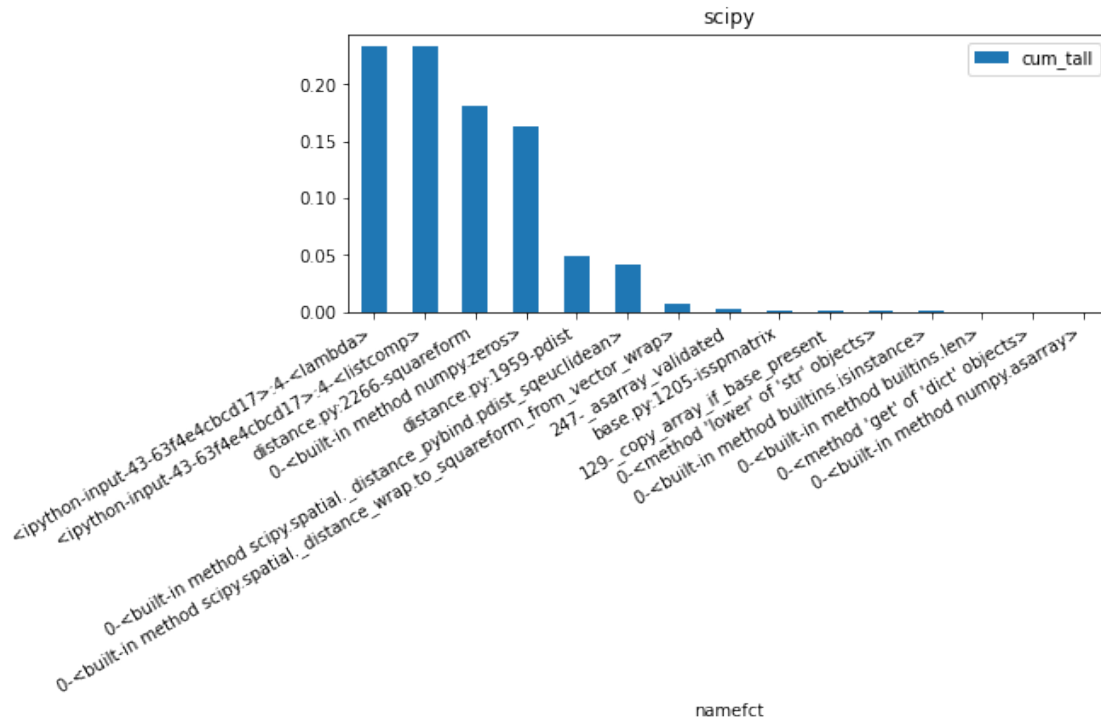


Curves are not linear and rather difficult to interpret. The algorithm *numpy-lower* and *scipy* should be close as the cost of both algorithm are similar. However, *scipy* reduces the number of trips between C and python. The C implementation of the distance is here: [sqeuclidean\\_distance\\_double](#). The final cost is a combination of computation, multithreading, allocations...

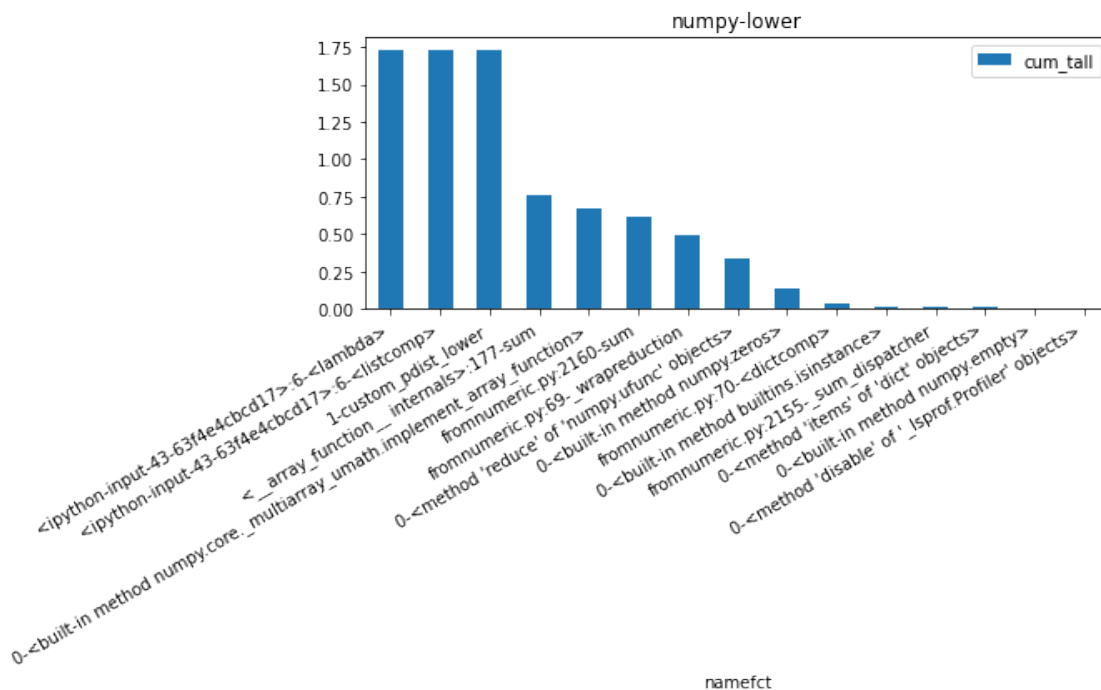
```
[30]: from pyquickhelper.pycode.profiling import profile
M = numpy.random.rand(100, 10)

pr1, df1 = profile(lambda: [squareform(pdist(M, metric='sqeuclidean')) for i in
    range(0, 1000)],
    as_df=True)
pr2, df2 = profile(lambda: [custom_pdist_lower(M) for i in range(0, 1000)], as_df=True)
```

```
[31]: ax = df1[['namefct', 'cum_tall']].head(n=15).set_index('namefct').plot(
    kind='bar', figsize=(8, 3), rot=30)
ax.set_title("scipy")
for la in ax.get_xticklabels():
    la.set_horizontalalignment('right')
```

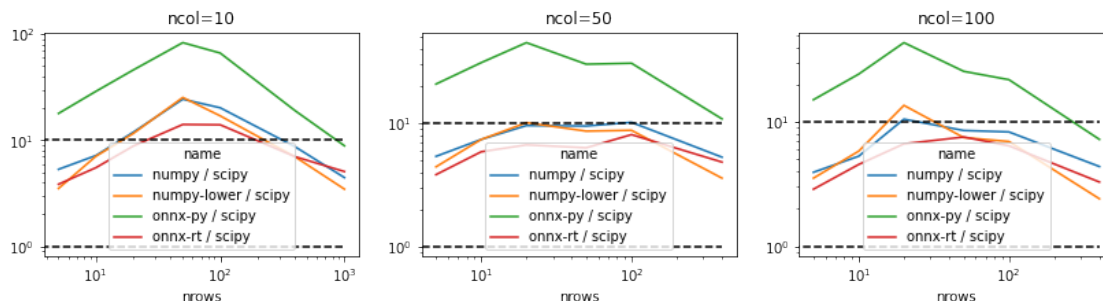


```
[32]: ax = df2[['namefct', 'cum_tall']].head(n=15).set_index('namefct').plot(
        kind='bar', figsize=(8, 3), rot=30)
ax.set_title("numpy-lower")
for la in ax.get_xticklabels():
    la.set_horizontalalignment('right');
```



Universal function do not seem to be very efficient in our case. The last graph shows time ratio between implementations of *pdist* and the baseline *scipy*.

```
[33]: fig, ax = plt.subplots(1, 3, figsize=(14, 3))
for i, ncol in enumerate([10, 50, 100]):
    piv = df[df.ncols==ncol].pivot("nrows", "name", "average")
    piv['numpy / scipy'] = piv['numpy'] / piv['scipy']
    piv['numpy-lower / scipy'] = piv['numpy-lower'] / piv['scipy']
    piv['onnx-py / scipy'] = piv['onnx-py'] / piv['scipy']
    piv['onnx-rt / scipy'] = piv['onnx-rt'] / piv['scipy']
    piv = piv[['numpy / scipy', 'numpy-lower / scipy',
               'onnx-py / scipy', 'onnx-rt / scipy']]
    piv.plot(ax=ax[i], logy=True, logx=True)
    ax[i].plot([0, max(piv.index)], [1, 1], '--', color='black')
    ax[i].plot([0, max(piv.index)], [10, 10], '--', color='black')
    ax[i].set_title("ncol=%d" % ncol)
ax;
```



## 1.5 Test with a new operator CDist

The final question is: *should we introduce a new operator into ONNX specifications?* The function *pdist* is not necessarily often used for a big number of observations as the square matrix it produces will even bigger. It seems reasonable. We showed that a python runtime based on *numpy* would not help, the implementation must be done in C++ or directly used the *scipy* version. The experiment was done with a *GaussianProcessRegressor*. The following section tests with and without a new operator *CDist* reusing *scipy* implementation.

```
[34]: import numpy
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import ExpSineSquared
from mlpredict.onnx_conv import to_onnx
from mlpredict.onnxrt import OnnxInference

iris = load_iris()
X, y = iris.data, iris.target
X_train, X_test, y_train, __ = train_test_split(X, y, random_state=12)
```

```

clr = GaussianProcessRegressor(ExpSineSquared(), alpha=20.)
clr.fit(X_train, y_train)

model_def = to_onnx(clr, X_train)

%onnxview model_def -r 1

```

[34]: <jyquickhelper.jspy.render\_nb\_js\_dot.RenderJsDot at 0x234ede21a00>

```

[35]: model_def_cdist = to_onnx(clr, X_train,
                                options={GaussianProcessRegressor: {'optim': 'cdist'}})
%onnxview model_def_cdist

```

[35]: <jyquickhelper.jspy.render\_nb\_js\_dot.RenderJsDot at 0x234edc44fd0>

```

[36]: oinf = OnnxInference(model_def)
oinf_cdist = OnnxInference(model_def_cdist)

```

```

[37]: %timeit oinf.run({'X': X_test})

```

4.24 ms ± 274 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

```

[38]: %timeit oinf_cdist.run({'X': X_test})

```

414 µs ± 15 µs per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```

[39]: oinfrt = OnnxInference(model_def, runtime="onnxruntime")
oinf_cdist = OnnxInference(model_def_cdist)

```

```

[40]: %timeit oinf_cdist.run({'X': X_test})

```

345 µs ± 26.8 µs per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

It is 10 times faster for this dataset so it is worth it. For bigger datasets, we should expect a lower gain but still significant.

[41]: